

000 001 002 003 004 005 006 007 008 009 010 011 012 013 014 015 016 017 018 019 020 021 022 023 024 025 026 027 028 029 030 031 032 033 034 035 036 037 038 039 040 041 042 043 044 045 046 047 048 049 050 051 052 053 TO TACKLE ADVERSARIAL TRANSFERABILITY: A NOVEL ENSEMBLE TRAINING METHOD WITH FOURIER TRANSFORMATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Ensemble methods are commonly used for enhancing robustness in machine learning. However, due to the “transferability” of adversarial examples, the performance of an ensemble model can be seriously affected even it contains a set of independently trained sub-models. To address this issue, we propose an efficient data transformation method based on a acute “weakness allocation” strategy, to diversify non-robust features. Our approach relies on a fine-grained analysis on the relation between non-robust features and adversarial attack directions. Moreover, our approach enjoys several other advantages, e.g., it does not require any communication between sub-models and the construction complexity is also quite low. We conduct a set of experiments to evaluate the performance of our proposed method and compare it with several popular baselines. The results suggest that our approach can achieve significantly improved robust accuracy over most existing ensemble methods, and meanwhile preserve high clean accuracy.

1 INTRODUCTION

In the past decade, *Deep neural networks (DNNs)* have achieved prominent performance on a broad range of real-world tasks (Goodfellow et al., 2016). However, a number of previous works show that DNNs are susceptible to carefully-crafted manipulations, where the manipulated data are called “adversarial examples” (Szegedy et al., 2014; Zhou et al., 2018; Heaven, 2019). The existence of adversarial examples severely impedes the application of DNNs in security-conscious scenarios, such as self-driving car (Rossolini et al., 2023; Zhu et al., 2021) and heath care (Newaz et al., 2020).

The *adversarial training* approach (Wang et al., 2023a; Madry et al., 2017) has gained significant attention due to its great effectiveness for defending against adversarial examples. However, the adversarial training approach often necessitates considerably high training time and large training dataset (Gowal et al., 2021; Carmon et al., 2019). Moreover, it has been observed that adversarial training is likely to incur certain decline in the accuracy on clean data, which also hinders the trained model to be applied for many practical tasks (Tsipras et al., 2018; Zhang et al., 2019).

Another important approach to enhance adversarial robustness is *ensemble training* (Tramèr et al., 2018). But recent studies (Demontis et al., 2019; Gao et al., 2022; Waseda et al., 2023) demonstrated that an adversarial example can attack different models even they are trained independently, and this phenomenon is the so-called “**transferability**” of adversarial examples. Hence, the strategy that simply integrates different models trained on the same original dataset is not sufficient to guarantee the overall robustness. To resolve this issue, different approaches have been proposed for maximizing the “diversity” among sub-models; in general, these approaches can be categorized into two classes: “simultaneous training” and “individual training” (Pang et al., 2019).

To reduce the similarity among sub-models, most existing “simultaneous training” methods attempt to incorporate some penalty during each epoch of parameter updates. Kariyappa & Qureshi (2019) proposed the “Gradient Alignment Loss (GAL)” method to minimize the gradient similarity between sub-models directly. Further, Yang et al. (2021) proposed the “Transferability Reduced Smooth (TRS)” method to improve GAL by adding a regularization term to increase the smoothness, as the models with a smoother loss function can reduce the “transferability” of attacks. Yang et al. (2020) aimed to isolate the adversarial vulnerability in each sub-model by distilling non-robust features, where the

054 sub-models can then generate diverse outputs being resilient against transfer attacks. Despite their
 055 effectiveness for defending adversarial attacks, the simultaneous training methods often require a
 056 substantial amount of memory since all the sub-models need to be stored in the GPUs in the training
 057 stage, which could be prohibitive if the number of sub-models is not small (say, more than 10) and/or
 058 their sizes are large. Additionally, the information interaction in parallel training can also cause extra
 059 large communication cost.

060 Different from simultaneous training, most “individual training” methods train each sub-model
 061 independently on a randomly transformed version of the given training dataset (Pang et al., 2019;
 062 AprilPyone & Kiya, 2021). This “random transformation” strategy yields diverse datasets, and thus
 063 different sub-models trained on these datasets can present diverse performances when confronting
 064 an adversarial attack. The individual training approach has higher flexibility and also requires less
 065 GPU memory, because the sub-models do not need to be stored simultaneously. Since there is
 066 no communication between sub-models, individual training methods are more suitable for parallel
 067 training with multiple GPUs. But unfortunately, recent studies showed that the commonly used
 068 random transformations (e.g. image cropping and rescaling) are not that effective under adversarial
 069 attacks (Athalye et al., 2018). The major cause of suppressing the performance of individual training
 070 is that the “transferability” problem is still not well addressed.

071 **Our contributions.** To tackle the transferability obstacle, we consider developing a new data
 072 transformation method for ensemble training. Our main contributions are summarized as follows:

- 073 – First, we propose a fine-grained analysis on the relation between non-robust features and adversarial
 074 attack directions (Section 3). Being different from the previous analysis on non-robust features, our
 075 new analysis provides us the hints that are particularly useful to allocate the potential vulnerability
 076 directions to a set of sub-models, and therefore paves the way for designing our ensemble training
 077 strategy.
- 078 – Second, we propose a data transform framework that can effectively promote the diversity of training
 079 data for robust ensemble training. The framework consists of two steps: “frequency selection” and
 080 “frequency transformation”, where the frequency is based on the Fourier transformation on the images.
 081 We propose two efficient frequency transformations with low complexities on the identified non-
 082 robust features. The first one is based on simple random noise, and the second one is a cute “targeted
 083 attack transformation” that can modify the non-robust features more effectively (Section 4.2).
- 084 – Finally, we conduct a set of experiments to evaluate the adversarial robustness of our approach on
 085 several benchmark datasets under the widely used attack algorithms. We also compare our approach
 086 with several open-source ensemble methods, such as ADP (Pang et al., 2019), GAL (Kariyappa &
 087 Qureshi, 2019), DVERGE (Yang et al., 2020), and TRS (Yang et al., 2021). Compared with those
 088 baselines, the experimental results suggest that our proposed approach can significantly outperform
 089 most of them in robust accuracy and also preserve comparable high clean accuracy.

090 1.1 OTHER RELATED WORKS

092 **Data transformation for ensemble training.** Guo et al. (2017) and Raff et al. (2019) proposed the
 093 transformations that preserve semantic information to reduce the impact of adversarial perturbation.
 094 AprilPyone & Kiya (2021) developed a training method that employs block-wise data transformations,
 095 where the input image is partitioned into blocks based on some private key. LINAC Rusu et al. (2022)
 096 uses a predetermined random seed (private key) to initialize and train a DNN to encode the input data,
 097 serving as an encrypted input transformation.

098 **Adversarial attack from frequency perspective.** Wang et al. (2020) explained that the model’s
 099 vulnerability to small distortions may be due to its dependence on high-frequency features. Yucel
 100 et al. (2023) proposed a data augmentation method that reduces the reliance on high-frequency
 101 components, so as to improve model’s robustness while maintaining clean accuracy. Maiya et al.
 102 (2021) and Bernhard et al. (2021) respectively showed that to fully understand the vulnerability, we
 103 should consider the distribution of the entire dataset with high and low frequencies.

104 2 PRELIMINARIES

105 **Some notations.** We consider the k -classification task: $\mathcal{X} \rightarrow \mathcal{Y}$ where \mathcal{X} is the input data space
 106 and $\mathcal{Y} = \{1, 2, \dots, k\}$ is the set of labels. A soft-classification model $f(\cdot; \beta)$ maps each $x \in \mathcal{X}$

to a vector $f(x; \beta) \in \mathbb{R}^k$, where β is the parameter vector that needs to be trained. Its associated hard-classification model is $F(\mathbf{x}; \beta) = \arg \max_i [f(x; \beta)]_i$, where $[\cdot]_i$ stands for the i -th coordinate. The model f is usually equipped with a loss function $\ell(f(\mathbf{x}; \beta), y)$, $x \in \mathcal{X}$ and $y \in \mathcal{Y}$, which is differentiable on β (e.g., cross-entropy loss). We refer to the accuracy on the original dataset as “**clean accuracy**” and the accuracy on adversarial examples as “**robust accuracy**”. We denote the one-hot k -dimensional vector that corresponds to the target label y as \mathbf{y}_{oh} .

Definition 2.1 (Ensemble Model) Let $\mathcal{M} = \{f_1, \dots, f_M\}$ be a set of sub-models for a k -classification task. We build the ensemble model with the following function:

$$f_{\text{E}}(x; \beta_{[1:M]}) = \frac{1}{M} \sum_{m \in [M]} \hat{F}_m(x; \beta_m), \quad (1)$$

where $\beta_{[1:M]} = \{\beta_m \mid 1 \leq m \leq M\}$, and $\hat{F}_m(x; \beta_m)$ is the one-hot k -dimensional vector of the hard-classification model $F_m(x; \beta_m)$ of f_m .

Definition 2.2 (Adversarial Attack and Targeted Attack) Given a model $f(\cdot; \beta)$ and an input $(x, y) \in \mathcal{X} \times \mathcal{Y}$, the **adversarial attack** algorithm \mathcal{A} returns a perturbed data x' inside the l_p ball of radius $\epsilon > 0$, which maximizes the loss function $\ell(f(\cdot; \beta), \cdot)$, or minimizes the loss function $\ell(f(\cdot; \beta), y_t)$ if given a target label $y_t \neq y$. For the latter one, we say it is a “**targeted attack from y to y_t** ”. Usually we set $p = 2$ or $p = \infty$.

As mentioned in Section 1, because our proposed approach is based on Fourier transform, we introduce several necessary notations below. Given an image x of size $L \times N$, the corresponding two-dimensional discrete Fourier transform can be written as: for any $0 \leq u \leq L - 1$ and $0 \leq v \leq N - 1$,

$$\tilde{x}[u, v] = \sum_{s=0}^{L-1} \sum_{t=0}^{N-1} x[s, t] \cdot e^{-2\pi j (\frac{us}{L} + \frac{vt}{N})}, \quad (2)$$

where “ j ” denotes the imaginary unit, and “ $\tilde{x}[u, v]$ ” is the entry in the u -th column and v -th row of the Fourier matrix \tilde{x} (“ $x[s, t]$ ” is defined similarly for the original image x). The pixels of the image x form the **time domain**, and the entries of \tilde{x} form the **frequency domain**. For a frequency (u, v) , the **amplitude** is the absolute value $|\tilde{x}[u, v]|$. We call a frequency (u, v) as a **frequency feature**.

3 FINE-GRAINED ANALYSIS ON ENSEMBLE MODEL VULNERABILITY

The previous work (Ilyas et al., 2019) categorizes the features learned by a model into robust and non-robust features. It shows that adversarial vulnerability is a natural consequence of the presence of highly predictive but non-robust features. Moreover, different models trained on the same dataset often have similar non-robust features, and therefore an adversarial example usually exhibits the unpleasant “transferability” property among them. Several other works also presented detailed discussions on the impact of non-robust features (Benz et al., 2021; Springer et al., 2021). Following those studies, a natural idea for tackling the transferability issue is to ensure that the sub-models should have **diverse non-robust features**. In this section, we provide a fine-grained analysis on the vulnerability of ensemble models and then conclude two important hints for achieving this “diversity” goal.

The following definitions are inspired by (Ilyas et al., 2019). Note that different from the term “feature” used in their article, we use “feature extractor” instead in our paper, since “feature” will be particularly used for referring to image feature in time or frequency domain. A model f should be comprised of a set of different feature extractors.

Definition 3.1 (Useful feature extractor) For a given data distribution $\mathcal{D} = \mathcal{X} \times \mathcal{Y}$, a feature extractor $\theta : \mathcal{X} \rightarrow \mathbb{R}^k$ is **useful**, if we have

$$\mathbb{E}_{(x, y) \sim \mathcal{D}} [\mathbf{y}_{\text{oh}}^\top \theta(x)] > \frac{1}{k}. \quad (3)$$

Recall that \mathbf{y}_{oh} is the one-hot k -dimensional vector of the label y . So the inequality (3) implies that a useful feature extractor can contribute to the model’s correct prediction.

162 **Definition 3.2 (robust and non-robust feature extractor)** We use $\mathcal{A}(x)$ to denote the adversarial
 163 example of a data item x as described in Definition 2.2. Let θ be a useful feature extractor. **(1)** We
 164 say θ is **robust** if the following condition holds for any i ($1 \leq i \leq k$):

$$\mathbb{E}_{(x,y) \sim \mathcal{D}_i} [\theta(\mathcal{A}(x))]_i > \frac{1}{k}$$

167 where \mathcal{D}_i represents the i -th class data. We denote the set of these robust feature extractors as Θ_R .

169 **(2)** The remaining useful feature extractors are **non-robust**. We assign these non-robust extractors
 170 to $k(k - 1)$ sets: $\{\Theta_{i,j} \mid 1 \leq i \neq j \leq k\}$ as follows. Initially, all these $k(k - 1)$ sets are empty.
 171 Then we go through all the non-robust feature extractors. For each non-robust θ , there must exist at
 172 least an index “ i ” such that

$$\mathbb{E}_{(x,y) \sim \mathcal{D}_i} [\theta(\mathcal{A}(x))]_i \leq 1/k;$$

175 we let $j = \arg \max_s \mathbb{E}_{(x,y) \sim \mathcal{D}_i} [\theta(\mathcal{A}(x))]_s$ and assign θ to $\Theta_{i,j}$ (note that j should be not equal to
 176 i , or there are multiple indices achieving the maximum expectation and at least one is not equal to
 177 i , since otherwise $\sum_{s=1}^k \mathbb{E}_{(x,y) \sim \mathcal{D}_i} [\theta(\mathcal{A}(x))]_s$ will less than 1). Eventually, these $k(k - 1)$ sets are
 178 constructed, where each $\Theta_{i,j}$ contains the feature extractors that are not robust to the attack from i to
 179 j .

180 **Remark 3.3** Intuitively, if a feature extractor is robust, it should have the capability to preserve
 181 its contribution to the correct prediction even under perturbation. It is also worth noting that a
 182 non-robust feature extractor θ could be assigned to multiple $\Theta_{i,j}$ s.

184 Assume we have a standardly trained model f consisting of a set of useful feature extractors, and we
 185 denote it as Θ_f . Each of them can be classified as robust or non-robust as Definition 3.2. Similar
 186 with the formulation proposed in (Ilyas et al., 2019), we can represent the model as

$$f(x) = \sum_{\theta \in \Theta_R \cap \Theta_f} w_\theta \theta(x) + \sum_{i,j=1, i \neq j}^k \sum_{\theta \in \Theta_{i,j} \cap \Theta_f} w_\theta \theta(x), \quad (4)$$

191 where each θ has a coefficient $w_\theta \in \mathbb{R}$. We then conduct our analysis based on Equation (4). Some
 192 recent works reveal that adversarial training method can obtain robust model through reducing the
 193 dependence on non-robust feature extractors (Allen-Zhu & Li, 2022; Tsipras et al., 2018). However,
 194 this strategy may cause certain downgrade performance on clean accuracy (because the non-robust
 195 feature extractors also contribute to obtaining correct prediction). Fortunately, we are able to avoid
 196 this dilemma in the context of ensemble training. **Namely, we just need to keep the non-robust**
 197 **features as diverse as possible, instead of entirely eliminating the dependence on those non-**
 198 **robust feature extractors.** To pave the way for realizing this goal, we introduce the definition of
 199 vulnerability of ensemble model bellow.

200 **Definition 3.4 (Vulnerability of ensemble model)** Suppose f_E is an ensemble model as described
 201 in Definition 2.1, and its associated hard-classification model is denoted by F_E : $\forall x$, $F_E(x) =$
 202 $\arg \max_i [f_E(x)]_i$. Given the data distribution $\mathcal{D} = \mathcal{X} \times \mathcal{Y}$, **the vulnerability of F_E** is defined as:

$$\text{Vr}(F_E) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) \neq y\}], \quad (5)$$

205 where $\mathbb{I}(\cdot)$ represents the indicator function. Furthermore, for any target class y_t , we can define **the**
 206 **vulnerability towards y_t** as $\text{Vr}(F_E, y_t) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\}]$.

208 The vulnerability of Definition 3.4 describes the success probability of an attack \mathcal{A} to the ensemble
 209 model F_E . We have the following key inequality, which indicates that $\text{Vr}(F_E)$ is bounded by
 210 considering all the attack directions, i.e.,

$$\text{Vr}(F_E) \leq \sum_{y_t \in \mathcal{Y}} \text{Vr}(F_E, y_t). \quad (6)$$

212 The proof of Inequality (6) is placed in Appendix A.1. Moreover, if $F_E(\mathcal{A}(x)) = y_t$, there are
 213 at least M/k sub-models returning the wrong label y_t due to the pigeonhole principle. Namely,

“ $\sum_{m=1}^M \mathbb{I}([f_m(\mathcal{A}(x))]_y < [f_m(\mathcal{A}(x))]_{y_t}) > \frac{M}{k}$ ” should be a necessary condition for successfully attacking from y to y_t . So it implies

$$\text{Vr}(F_E, y_t) \leq \mathbb{E}_{(x,y) \sim \mathcal{D}} \left[\mathbb{I} \left(\sum_{m=1}^M \mathbb{I}([f_m(\mathcal{A}(x))]_y < [f_m(\mathcal{A}(x))]_{y_t}) > \frac{M}{k} \right) \right]. \quad (7)$$

From the upper bound (6), we can decrease the total vulnerability by reducing $\text{Vr}(F_E, y_t)$ for each y_t . Also, from (7) we know that $\text{Vr}(F_E, y_t)$ can be reduced by decreasing the chance of “[$f_m(\mathcal{A}(x))]$ _y < [$f_m(\mathcal{A}(x))]$ _ y_t ” over $m \in \{1, 2, \dots, M\}$. According to the Equation (4), the inequality “[$f_m(\mathcal{A}(x))]$ _y < [$f_m(\mathcal{A}(x))]$ _ y_t ” can be rewritten as

$$\begin{aligned} & \left[\sum_{\theta \in \Theta_R^m} w_\theta \theta(\mathcal{A}(x)) + \sum_{i,j=1, i \neq j}^k \sum_{\theta \in \Theta_{i,j}^m} w_\theta \theta(\mathcal{A}(x)) \right]_y < \\ & \left[\sum_{\theta \in \Theta_R^m} w_\theta \theta(\mathcal{A}(x)) + \sum_{i,j=1, i \neq j}^k \sum_{\theta \in \Theta_{i,j}^m} w_\theta \theta(\mathcal{A}(x)) \right]_{y_t}, \end{aligned} \quad (8)$$

where Θ_R^m and $\Theta_{i,j}^m$ respectively denote the sets of robust and non-robust feature extractors for the sub-model f_m . Moreover, the set Θ_{y,y_t}^m should have relatively larger influence to the value of (8) than other feature extractor set $\Theta_{y,j}^m$ with $j \neq y_t$, due to the outer operator “[.] $_{y_t}$ ”. Therefore, we conclude our first hint as an intuition for reducing $\text{Vr}(F_E)$.

Hint (i): To decrease the vulnerability in the attack direction y_t (i.e., each term $\text{Vr}(F_E, y_t)$ in the upper bound of (6)), it is reasonable to decrease the influence from the non-robust feature extractors of Θ_{y,y_t}^m .

In Hint (i), a major difference from the previous analysis (Ilyas et al., 2019; Allen-Zhu & Li, 2022) is that, we in particular relate each attack direction y_t to some specific non-robust feature extractors, where the benefit is that these correspondences can effectively help us to build the diverse ensemble model. Moreover, According to the principle of ensemble methods, as long as at least $M/2 + 1$ sub-models are not successfully attacked, the ensemble model will successfully defend against the attack. So we conclude the second hint that is also important for designing our approach.

Hint (ii): For each attack direction y_t , we only need to consider manipulating the training data of $M/2 + 1$ sub-models instead of all the M sub-models.

Overall, the above Hint (i) & (ii) play the key roles for inspiring our data transformation method in Section 4.

4 OUR ENSEMBLE TRAINING METHOD

We first introduce our model and high-level idea in Section 4.1, and then elaborate on the technical details for the data transformations in Section 4.2.

4.1 OVERVIEW OF OUR FRAMEWORK

Note that the feature extractors of a model depend on the given training data. Namely, any modification on the features of the training data can implicitly influence the model. Thus, in this section we follow the Hint (i) & (ii) of Section 3 to design an effective data transformation method. The transformation is expected to modify the features of the training data, so as to enhance the robustness of the trained ensemble model. We train a set of distinct sub-models on the transformed training data; these sub-models can be integrated into an ensemble model being robust against adversarial attacks, while preserving the clean accuracy of each sub-model as much as possible. We use “ π_m ” to denote the transformation for the m -th sub-model, $1 \leq m \leq M$, and formulate the following problem by slightly modifying Definition 2.1 (replace x by the adversarial example $\mathcal{A}(x)$ for each sub-model):

$$\min \mathbb{E}_{(x,y) \sim \mathcal{X} \times \mathcal{Y}} \ell \left(\frac{1}{M} \sum_{m \in [M]} \hat{F}_m(\mathcal{A}(x); \beta_m), y \right) \quad (9)$$

where β_m is obtained by training on the transformed data, i.e., $\beta_m = \arg \min_{\beta} \mathbb{E}_{(x,y) \sim \mathcal{X} \times \mathcal{Y}} \ell(f_m(\pi_m(x); \beta), y)$ for each $m \in \{1, 2, \dots, M\}$.

The major challenge for solving the above problem (9) is how to design a set of appropriate transformations $\{\pi_m \mid 1 \leq m \leq M\}$, so that the obtained parameters $\beta_{[1:M]}$ can yield sufficiently diverse sub-models. To address this issue, we leverage the transformation from frequency domain to guide the non-robust features of each sub-model to be as diverse as possible. Specifically, we introduce a method called “**Frequency Domain Transformation (FDT)**” for constructing the set of diverse training datasets $\{\pi_1(\mathcal{X}), \pi_2(\mathcal{X}), \dots, \pi_M(\mathcal{X})\}$. FDT relies on a key “weakness allocation” strategy. Roughly speaking, the strategy aims to promote the diversity of the constructed datasets, and meanwhile preserve that the overall clean accuracy should not be sacrificed in the ensemble. The details are presented in the next section.

4.2 FREQUENCY DOMAIN TRANSFORMATION

Before performing the transformation, we need to select a set of non-robust features. In time domain, a simple observation is that an image feature is usually invariant under spatial translation, e.g., it can appear at different positions in images. This property causes the challenge for directly identifying and representing non-robust features in time domain. Thus we turn our attention to the frequency domain. Moreover, some previous studies on robust learning already revealed that robust and non-robust features are often deeply related to frequency domain (Wang et al., 2020; Bernhard et al., 2021; Maiya et al., 2021).

Amplitude based selection. To identify the non-robust frequencies, a straightforward idea is to test the robustness of each individual frequency and select the non-robust ones. Nevertheless, it may takes a large computational cost since the number of frequencies is high (e.g., if the input image is 64×64 , the number of frequencies is also $64 \times 64 \approx 4 \times 10^3$). We propose an easy-to-implement selection idea based on the amplitudes, since the amplitudes can be directly obtained via the Fourier transformation with low complexity. According to the previous research (Ilyas et al., 2019; Benz et al., 2021; Springer et al., 2021), a feature can be regarded as “robust” if it cannot be easily manipulated by small perturbations. We observe that high-amplitude frequency features usually dominate the ground truth of an image. Figure 4 in our Appendix illustrates an example to show that, if we keep high-amplitude frequencies and remove low-amplitude ones, the image is changed slightly even with adding certain noise (i.e., we can still recognize the ground truth from the modified image). This observation suggests that high-amplitude frequency features are more strongly related to the semantic information of image. So in our following approach, we maintain high-amplitude frequency features as “robust features”, and select the frequencies with low amplitudes (by setting a threshold “ τ ”) to transform. Moreover, we can conveniently observe the performance changing through varying the threshold τ in our experiment. Figure 1 illustrates the amplitude-based selection.

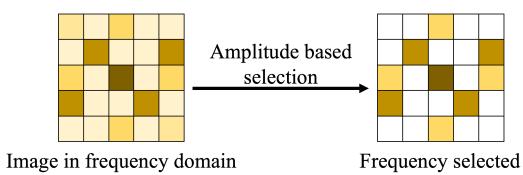


Figure 1: We use a 5×5 image as a toy example, where the intensity of the color indicates the magnitude of the amplitude. In our amplitude-based selection, we retain the high-amplitude frequencies (i.e., the darker regions) and perform data transformations on the low-amplitude frequencies (i.e., the white regions).

ing the images to different targeted classes, and then use them to replace the selected non-robust frequencies.

Targeted-attack transformation: We briefly explain our intuition first. It was shown that adversarial attacks have the capability to manipulate non-robust features (Ilyas et al., 2019; Yang et al., 2020). In particular, a targeted attack as introduced in Definition 2.2, aims at modifying non-robust features that are associated with a specific target label. For instance, let us consider a data point (x, y) in the original dataset $\mathcal{X} \times \mathcal{Y}$; we set the target label as y_t and obtain the corresponding adversarial example

Following our frequency selection, we propose two transformation methods for promoting the diversity by using the identified non-robust features. Our first approach is from a straightforward idea, which is just to replace the non-robust features by random noise (due to the space limit, we leave the details to Appendix C). This method is very easy to implement in practice. Though it can achieve certain degree of improvement upon previous ensemble training methods, the performance is still not very promising (as shown in our experiments). To further improve the effectiveness, we propose a more sophisticated approach called “targeted-attack transformation”, which constructs a set of different “substitute” features through attack-

324 x' (x' contains the modified non-robust features that are associated with y_t). When training a model
 325 using (x', y) , intuitively it can be viewed as an “immunization” for defending the attack from y to y_t ;
 326 and consequently, the chance that obtaining the wrong label y_t for the data with label y decreases. In
 327 other words, it becomes more difficult to attack the images with label y to y_t than to the other classes.
 328 We call the modified non-robust feature as a “substitute” feature derived by the targeted attack.

329 Motivated by this observation, we can construct different transformations by using $k \times (k - 1)$
 330 targeted attacks (since each label can be attacked to be the other $k - 1$ labels); these attacks can yield
 331 different substitute features, and then we use these features to replace the corresponding non-robust
 332 features in the original dataset (based on Hint (i) in Section 3); finally, the M transformed datasets are
 333 obtained via an allocation algorithm, where each substitute feature is captured by at least $M/2 + 1$
 334 datasets (based on Hint (ii) in Section 3). Overall, due to the completeness of the $k \times (k - 1)$
 335 targeted attacks, the M sub-models trained on those datasets can guarantee the robustness of the final
 336 ensemble solution. We introduce some definitions for our transformation first.

337 **Definition 4.1 (Strengthen a dataset)** *Let $y_1 \neq y_2 \in \mathcal{Y}$. If a given training dataset P contains at
 338 least one adversarial example who has the original label y_1 but is misclassified as y_2 , we say that P
 339 has been strengthened by the attack direction from y_1 to y_2 (“ $\overrightarrow{y_1 y_2}$ -direction” for short).*

341 In other words, if P is not strengthened in $\overrightarrow{y_1 y_2}$ -direction, the model trained on P is more likely to
 342 be fragile to the targeted attacks from y_1 to y_2 . Also, the dataset P may have not been strengthened
 343 in multiple different directions. So we define its “weakness set” $\mathcal{W} = \{\overrightarrow{y_1 y_2} \mid 1 \leq y_1, y_2 \leq k, y_1 \neq$
 344 $y_2\}$, and P has not been strengthened in $\overrightarrow{y_1 y_2}$ -direction}.

345 **Definition 4.2 (Diversity of weakness sets)** *Given M datasets $\{P_1, P_2, \dots, P_M\}$ with their corre-
 346 sponding weakness sets $\{\mathcal{W}_1, \mathcal{W}_2, \dots, \mathcal{W}_M\}$, we define their diversity:*

$$\text{Div}(P_1, P_2, \dots, P_M) = 1 - \frac{|\mathcal{W}_1 \cap \mathcal{W}_2 \cap \dots \cap \mathcal{W}_M|}{\max\{|\mathcal{W}_1|, |\mathcal{W}_2|, \dots, |\mathcal{W}_M|\}}.$$

351 It is easy to see that the higher the value $\text{Div}(P_1, P_2, \dots, P_M)$, the more diverse the corresponding
 352 weakness sets. A higher diversity suggests that the vulnerabilities of the M sub-models trained
 353 on those datasets are more likely to be different. To achieve a nice performance in terms of both
 354 accuracy and robustness, we need to take account of the diversity function “ Div ” for designing the
 355 transformations. The basic principle is:

356 *On the one hand, our transformed datasets should have a sufficiently large number of diverse
 357 substitute features, so that one adversarial attack cannot easily capture more than half of the M
 358 sub-models. On the other hand, the datasets should also maintain the major information of the
 359 original input as much as possible, since otherwise the clean accuracy may decline due to the added
 360 substitute features.*

361 To provide an appropriate trade-off, we propose the following constrained optimization objective: let
 362 \mathbb{C} be the set of all the $\binom{M}{M/2}$ combinations of $M/2$ -size subsets from $\{1, 2, \dots, M\}$, and then

$$\max_{P_1, P_2, \dots, P_M} \min\{|\mathcal{W}_1|, |\mathcal{W}_2|, \dots, |\mathcal{W}_M|\} \quad (10)$$

$$\text{s.t. } \forall \{i_1, i_2, \dots, i_{M/2}\} \in \mathbb{C}, \quad \text{Div}(P_{i_1}, P_{i_2}, \dots, P_{i_{M/2}}) = 1. \quad (11)$$

367 We maximize the objective function of (10) because we want to minimize the modification degree
 368 for each transformed dataset. Intuitively, a large weakness set indicates that the corresponding
 369 dataset is not changed significantly by the transformation, and thus the clean accuracy is likely
 370 to be well preserved. The constraint (11) guarantees that any $M/2$ datasets have the intersection
 371 $\mathcal{W}_{i_1} \cap \mathcal{W}_{i_2} \cap \dots \cap \mathcal{W}_{i_{M/2}} = \emptyset$, that is, they do not share any common direction. Consequently, the
 372 ensemble solution should be robust to any attack direction. To achieve this twofold goal, we design
 373 an efficient allocation strategy together with an attack-guided transformation on the training data.
 374 Specifically, the procedure consists of the following two stages.

375 **Stage (1): allocating the weakness sets to the sub-models.** For each $\overrightarrow{y_1 y_2}$ -direction, $1 \leq y_1, y_2 \leq$
 376 k , there are at most $\lceil \frac{M}{2} \rceil - 1$ sets that contain this direction (due to the constraint (11)), so the sum
 377 $\sum_{1 \leq i \leq M} |\mathcal{W}_i|$ is no larger than $k(k - 1) * (\lceil \frac{M}{2} \rceil - 1)$. Therefore, the maximum value of Eq (10) is
 no larger than

$$k(k-1) * (\lceil \frac{M}{2} \rceil - 1)/M \quad (12)$$

based on the pigeonhole principle. We assign the total $k(k-1) * (\lceil \frac{M}{2} \rceil - 1)$ directions (each direction is duplicated to be $\lceil \frac{M}{2} \rceil - 1$ copies) to M sets in a round-robin way, where the number of directions assigned to each set is no larger than the upper bound (12). Please refer to Figure 2 for an example.

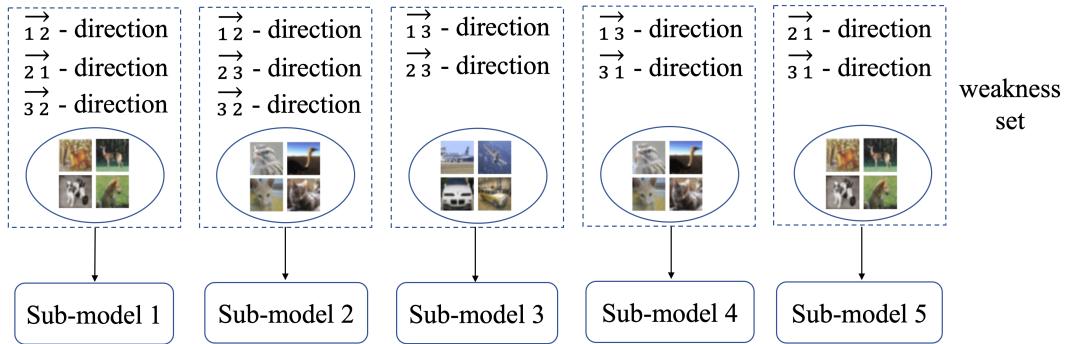


Figure 2: Assign the attack directions to five sub-models for a three-class classification task.

Stage (2): constructing the new datasets. Following the allocation, we transform the original dataset, denoted by P_{ori} , to align with the assigned weakness sets for the M sub-models correspondingly. Using the same notations of Definition 4.2, we denote the waiting-to-construct dataset for the m -th sub-model as P_m (which is initialized to be \emptyset), $1 \leq m \leq M$. First, we divide P_{ori} into k subsets C_1, C_2, \dots, C_k , where each C_j corresponds to the label j , for $1 \leq j \leq k$; further, each C_j is equally partitioned to $k-1$ disjoint parts $\{C_{j,1}, C_{j,2}, \dots, C_{j,k-1}\}$ at random. For each data (x, j) in $C_{j,i}$, we attack it from j to h (let $h = i + j \bmod k$) to obtain the adversarial perturbation; then we only substitute the low-amplitude frequencies of x with the perturbation, and other frequencies (which have their amplitudes higher than the aforementioned threshold τ) remain unchanged. We denote the new dataset as $C'_{j,i}$. Finally, we add $C'_{j,i}$ to P_m if the $i \rightarrow h$ -direction is not in the weakness set \mathcal{W}_m . From the construction method of the weakness sets, we know that the $i \rightarrow h$ -direction can appear in at most $\lceil \frac{M}{2} \rceil - 1$ weakness sets. So, the set $C'_{j,i}$ can be added to at least $\lceil \frac{M}{2} \rceil$ different P_m s. Consequently, the completeness for defending the $k(k-1)$ attack directions can be guaranteed, i.e., the constraint (11) is satisfied. Due to the space limit, the full details are shown in Appendix D. Figure 3 shows the schematic diagram of the construction process, and the full details are shown in Appendix D.

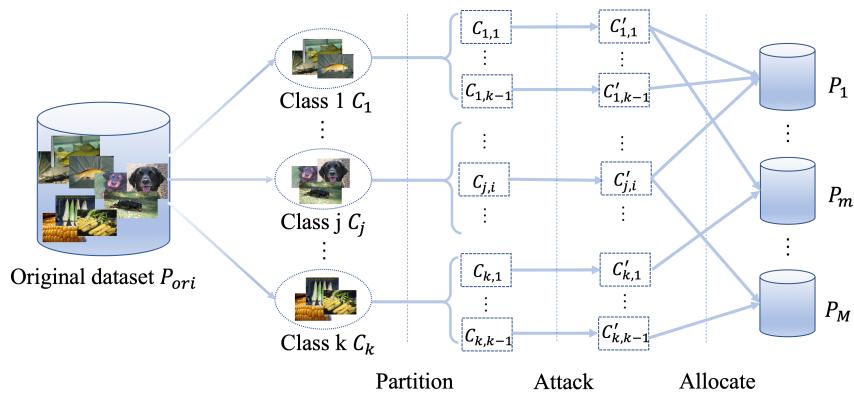


Figure 3: A schematic diagram of the construction process. In the allocation stage, each $C'_{j,i}$ is added to P_m if the $i \rightarrow h$ -direction is not in the weakness set \mathcal{W}_m , $h = i + j \bmod k$.

Remark 4.3 We are aware of some previous robust learning approaches that also depend on data modification (Allen-Zhu & Li, 2022; Tsipras et al., 2018). But their approaches usually tend to completely eliminate non-robust features. Our method is quite different, where the goal is to leverage the carefully selected non-robust features to weaken the transferability among sub-models. For each sub-model, we only modify the non-robust features corresponding to certain directions, rather than all non-robust features, and therefore the modification yields relatively lower impact on clean accuracy. Moreover, we partition each class (C_j) into $k - 1$ subsets $\{C_{j,1}, C_{j,2}, \dots, C_{j,k-1}\}$, with each subset being attacked to a specified class. This step eliminates the need to attack each data point across all classes, thereby reducing the computational complexity of constructing the new datasets.

5 EXPERIMENTS

We conduct our experiments on the widely used image datasets CIFAR-10, CIFAR-100 (Krizhevsky & Hinton, 2009), and Tiny-ImageNet-200 (Deng et al., 2009). As for the baselines, we reproduce the existing ensemble models including ADP (Pang et al., 2019), GAL (Kariyappa & Qureshi, 2019), DVERGE (Yang et al., 2020), and TRS (Yang et al., 2021), with their released codes and recommended hyperparameter settings. As for our approach, “FDT-random” and “FDT-target” respectively denote the methods utilizing random noise based transformation and target-attack transformation; “FDT-hybrid” represents the method that combines both, that is, we set two frequency selection thresholds τ_1 and τ_2 ($\tau_1 < \tau_2$), and perform random and target-attack transformations on the frequencies less than τ_1 and the frequencies between τ_1 and τ_2 , respectively (due to the space limit, more details are shown in Appendix E). We train each sub-model based on ResNet-20 (He et al., 2016) and use Adam optimizer (Kingma & Ba, 2015) with an initial learning rate of 0.001 for 200 epochs. To further test their performance on neural network with larger scale, we also use WideResNet28-10 (Zagoruyko & Komodakis, 2016) to train the sub-models and the results are placed in our supplement. All the experiments are implemented with PyTorch (Paszke et al., 2017) on a single NVIDIA GeForce RTX 3090 with 24GB of memory and 1TB of storage. We assess the performance of our models through 5 repeated runs and compute error bars. Utilizing the numpy library, we calculate the standard deviation and subsequently derive the standard error of the mean (SEM).

Varying the number of sub-models. We take the ResNet-20 model trained on CIFAR-10 as an example and test the performance of FDT with different numbers of sub-models in the ensemble. In this experiment, we set the frequency selection threshold τ to be 0.8. Then we evaluate the performance of FDT-hybrid under FGSM (Madry et al., 2017) , PGD (Carlini et al., 2019), and AutoAttack (AA) (Croce & Hein, 2020) attack methods with l_∞ perturbations of size $\epsilon = 0.02$. The results in Table 1 indicate that our clean accuracy has relatively smaller change as the number increases, while the robust accuracy can be substantially improved from 3 to 20 sub-models.

Table 1: Performance of FDT-hybrid with different sub-model numbers on CIFAR-10.

Sub-model numbers	3	5	8	12	20
Clean accuracy	90.20 ± 0.03	90.75 ± 0.03	91.35 ± 0.05	91.51 ± 0.06	91.86 ± 0.07
FGSM ($\epsilon = 0.02$)	58.04 ± 0.13	61.66 ± 0.15	62.41 ± 0.11	63.96 ± 0.12	64.27 ± 0.14
PGD ($\epsilon = 0.02$)	20.01 ± 0.04	26.10 ± 0.07	29.20 ± 0.05	29.78 ± 0.08	29.71 ± 0.07
AutoAttack ($\epsilon = 0.02$)	19.42 ± 0.04	25.37 ± 0.05	27.33 ± 0.04	28.12 ± 0.07	28.92 ± 0.07

Results for white-box attack. To maintain consistency with the baseline ensemble methods from the literature, we ensemble three ResNet-20 sub-models here and evaluate the robust accuracy using $\epsilon = 0.01$ and $\epsilon = 0.02$. In the white-box attack setting, the attacker has full knowledge of the models, including model parameters, architecture, and ensemble training strategy. To evaluate the adversarial robustness of the ensemble, we conduct the following white-box attacks: PGD, FGSM, BIM (Goodfellow et al., 2015), MIM (Dong et al., 2018), C&W (Carlini & Wagner, 2017) and AutoAttack (AA). The attacks are implemented using AdverTorch (Ding et al., 2019). We take the robust and clean accuracies, and average training time per epoch as the evaluation metrics.

Table 2 presents the obtained robust accuracies of the baseline ensemble methods on CIFAR-10 and CIFAR-100. In addition, we show the average training time per epoch of different ensemble methods. The experimental results suggest that our FDT-random method can achieve higher adversarial

486
 487 Table 2: Robust and Clean Accuracy (%) and average training time of different ensemble methods
 488 against white-box attacks on CIFAR-10 and CIFAR-100. “ ϵ ” and “ λ ” stand for the l_∞ norm of the
 489 adversarial perturbation and the coefficient of C&W attack respectively.

CIFAR-10	ADP	GAL	DVERGE	TRS	FDT-random	FDT-target	FDT-hybrid
Clean accuracy	91.84	91.81	91.37	69.60	89.88 \pm 0.02	90.16 \pm 0.04	90.20 \pm 0.03
FGSM ($\epsilon=0.01$)	59.48	44.97	70.05	34.54	66.96 \pm 0.12	72.88 \pm 0.12	72.24 \pm 0.12
FGSM ($\epsilon=0.02$)	53.38	30.58	56.33	21.65	46.28 \pm 0.10	55.54 \pm 0.09	58.04 \pm 0.13
PGD ($\epsilon=0.01$)	14.45	1.35	40.55	13.52	45.42 \pm 0.09	46.58 \pm 0.07	48.48 \pm 0.09
PGD ($\epsilon=0.02$)	2.95	0.34	11.49	2.22	12.24 \pm 0.03	15.08 \pm 0.05	20.01 \pm 0.04
BIM ($\epsilon=0.01$)	14.15	1.37	40.51	13.71	45.24 \pm 0.03	46.86 \pm 0.04	48.57 \pm 0.05
BIM ($\epsilon=0.02$)	3.01	0.27	10.65	1.98	11.68 \pm 0.03	14.86 \pm 0.03	16.63 \pm 0.02
MIM ($\epsilon=0.01$)	20.38	2.05	44.74	15.85	47.73 \pm 0.05	49.97 \pm 0.06	51.48 \pm 0.07
MIM ($\epsilon=0.02$)	5.11	0.69	14.76	3.01	15.14 \pm 0.04	18.27 \pm 0.02	20.09 \pm 0.03
AA ($\epsilon=0.01$)	1.80	0.00	43.34	10.86	46.09 \pm 0.09	48.83 \pm 0.08	51.56 \pm 0.08
AA ($\epsilon=0.02$)	0.00	0.00	13.72	0.60	9.38 \pm 0.05	15.70 \pm 0.05	19.42 \pm 0.04
C&W ($\lambda=0.1$)	20.96	31.57	52.35	18.94	45.01 \pm 0.10	55.48 \pm 0.10	56.08 \pm 0.11
CIFAR-100	ADP	GAL	DVERGE	TRS	FDT-random	FDT-target	FDT-hybrid
Clean accuracy	67.04	67.70	66.16	66.63	66.29 \pm 0.11	67.64 \pm 0.08	66.70 \pm 0.09
FGSM ($\epsilon=0.01$)	17.82	16.89	33.94	26.22	35.42 \pm 0.12	40.46 \pm 0.14	39.85 \pm 0.14
FGSM ($\epsilon=0.02$)	10.53	7.80	26.61	18.99	22.40 \pm 0.05	32.30 \pm 0.06	30.27 \pm 0.08
PGD ($\epsilon=0.01$)	0.80	0.11	14.62	0.89	21.54 \pm 0.06	22.19 \pm 0.05	24.93 \pm 0.05
PGD ($\epsilon=0.02$)	0.01	0.02	4.25	0.08	4.84 \pm 0.03	7.27 \pm 0.02	8.63 \pm 0.03
BIM ($\epsilon=0.01$)	0.68	0.23	14.85	0.86	21.10 \pm 0.09	22.39 \pm 0.07	24.35 \pm 0.06
BIM ($\epsilon=0.02$)	0.02	0.0	4.07	0.09	4.80 \pm 0.04	6.80 \pm 0.05	8.40 \pm 0.05
MIM ($\epsilon=0.01$)	0.78	0.12	16.82	1.17	23.14 \pm 0.06	24.68 \pm 0.10	27.09 \pm 0.09
MIM ($\epsilon=0.02$)	0.01	0.02	5.31	0.08	6.47 \pm 0.03	8.87 \pm 0.04	10.19 \pm 0.05
AA ($\epsilon=0.01$)	0.01	0.00	11.23	0.07	16.02 \pm 0.09	16.03 \pm 0.09	16.41 \pm 0.12
AA ($\epsilon=0.02$)	0.00	0.00	2.72	0.04	3.12 \pm 0.04	4.54 \pm 0.05	5.47 \pm 0.07
C&W ($\lambda=0.1$)	0.74	3.70	10.68	7.77	25.07 \pm 0.10	29.43 \pm 0.09	30.66 \pm 0.13
Time (s)	ADP	GAL	DVERGE	TRS	FDT-random	FDT-target	FDT-hybrid
CIFAR-10	30.15	69.92	134.33	350.42	37.04	108.22	114.23
CIFAR-100	30.34	69.71	129.25	344.92	37.12	108.43	113.87

519
 520 robustness over other baselines on both CIFAR-10 and CIFAR-100, with the training time only higher
 521 than ADP (and much lower than other baselines). Furthermore, the FDT-hybrid ensemble method
 522 achieves an even better robustness than FDT-random, though its running time is higher since it needs
 523 to perform the target-attack transformation.

524 **Summary on other experimental results placed in Appendix F.** We also conduct the experiments
 525 to examine the performance of FDT under black-box attack, and assess the transferability of our
 526 method across various sub-models. The results indicate the competitive robustness of our method in
 527 defending against black-box attacks. Then, we evaluate the trade-off between clean accuracy and
 528 robust accuracy by varying the frequency selection threshold τ . The result shows that the ensemble
 529 model has lower clean accuracy and higher robust accuracy with the increasing of τ . Moreover, we
 530 included some ablation studies on datasets and model architectures. These experiments demonstrate
 531 that our method performs the best among ensemble-based baseline methods.

532 6 CONCLUSION AND FUTURE WORK

533 In this paper, we present a novel data transformation approach to improve the robustness of ensemble
 534 models against adversarial attacks. By leveraging the frequency based features and strategically
 535 allocating adversarial examples, we demonstrate the effectiveness of our method in enhancing
 536 adversarial robustness while maintaining high accuracy on clean data. As for the future work, we
 537 can consider other types of transformation methods (e.g., beyond using frequency) to improve the
 538 ensemble robustness. Also, it is interesting to consider more complicated scenarios for ensemble
 539 training, such as federated learning with concerning the privacy issue.

540 REFERENCES
541

- 542 Zeyuan Allen-Zhu and Yuanzhi Li. Feature purification: How adversarial training performs robust
543 deep learning. In *2021 IEEE 62nd Annual Symposium on Foundations of Computer Science (FOCS)*, pp. 977–988. IEEE, 2022.
- 544
- 545 MaungMaung AprilPyone and Hitoshi Kiya. Block-wise image transformation with secret key
546 for adversarially robust defense. *IEEE Transactions on Information Forensics and Security*, 16:
547 2709–2723, 2021.
- 548
- 549 Anish Athalye, Nicholas Carlini, and David Wagner. Obfuscated gradients give a false sense of
550 security: Circumventing defenses to adversarial examples. In *International conference on machine
551 learning*, pp. 274–283. PMLR, 2018.
- 552 Philipp Benz, Chaoning Zhang, and In So Kweon. Batch normalization increases adversarial
553 vulnerability and decreases adversarial transferability: A non-robust feature perspective. In
554 *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 7818–7827,
555 2021.
- 556 Rémi Bernhard, Pierre-Alain Moëlllic, Martial Mermilliod, Yannick Bourrier, Romain Cohendet,
557 Miguel Solinas, and Marina Reyboz. Impact of spatial frequency based constraints on adversarial
558 robustness. In *2021 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8. IEEE,
559 2021.
- 560
- 561 Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks. In *2017
562 ieee symposium on security and privacy (sp)*, pp. 39–57. Ieee, 2017.
- 563
- 564 Nicholas Carlini, Anish Athalye, Nicolas Papernot, Wieland Brendel, Jonas Rauber, Dimitris Tsipras,
565 Ian Goodfellow, Aleksander Madry, and Alexey Kurakin. On evaluating adversarial robustness.
566 *arXiv preprint arXiv:1902.06705*, 2019.
- 567
- 568 Yair Carmon, Aditi Raghunathan, Ludwig Schmidt, John C Duchi, and Percy S Liang. Unlabeled
569 data improves adversarial robustness. *Advances in Neural Information Processing Systems*, 32,
2019.
- 570
- 571 Francesco Croce and Matthias Hein. Reliable evaluation of adversarial robustness with an ensemble
572 of diverse parameter-free attacks. In *International conference on machine learning*, pp. 2206–2216.
573 PMLR, 2020.
- 574 Ambra Demontis, Marco Melis, Maura Pintor, Jagielski Matthew, Battista Biggio, Oprea Alina, Nita-
575 Rotaru Cristina, Fabio Roli, et al. Why do adversarial attacks transfer? explaining transferability
576 of evasion and poisoning attacks. In *28th USENIX security symposium*, pp. 321–338. USENIX
577 Association, 2019.
- 578
- 579 Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale
580 hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*,
581 pp. 248–255. Ieee, 2009.
- 582
- 583 Gavin Weiguang Ding, Luyu Wang, and Xiaomeng Jin. Advertorch v0. 1: An adversarial robustness
toolbox based on pytorch. *arXiv preprint arXiv:1902.07623*, 2019.
- 584
- 585 Yingpeng Dong, Fangzhou Liao, Tianyu Pang, Hang Su, Jun Zhu, Xiaolin Hu, and Jianguo Li. Boosting
586 adversarial attacks with momentum. In *Proceedings of the IEEE conference on computer vision
587 and pattern recognition*, pp. 9185–9193, 2018.
- 588
- 589 Xitong Gao, Cheng-Zhong Xu, et al. Mora: Improving ensemble robustness evaluation with model
reweighing attack. *Advances in Neural Information Processing Systems*, 35:26955–26965, 2022.
- 590
- 591 Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial
592 examples. *stat*, 1050:20, 2015.
- 593
- Ian J. Goodfellow, Yoshua Bengio, and Aaron C. Courville. *Deep Learning*. Adaptive computation
and machine learning. MIT Press, 2016. ISBN 978-0-262-03561-3.

- 594 Sven Gowal, Sylvestre-Alvise Rebuffi, Olivia Wiles, Florian Stimberg, Dan Andrei Calian, and
 595 Timothy A Mann. Improving robustness using generated data. *Advances in Neural Information
 596 Processing Systems*, 34:4218–4233, 2021.
- 597 Chuan Guo, Mayank Rana, Moustapha Cisse, and Laurens Van Der Maaten. Countering adversarial
 598 images using input transformations. *arXiv preprint arXiv:1711.00117*, 2017.
- 600 Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image
 601 recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*,
 602 pp. 770–778, 2016.
- 603 Douglas Heaven. Why deep-learning ais are so easy to fool. *Nature*, 574(7777):163–166, 2019.
- 604 Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Logan Engstrom, Brandon Tran, and Aleksander
 605 Madry. Adversarial examples are not bugs, they are features. *Advances in neural information
 606 processing systems*, 32, 2019.
- 607 Sanjay Kariyappa and Moinuddin K Qureshi. Improving adversarial robustness of ensembles with
 608 diversity training. *arXiv e-prints*, pp. arXiv–1901, 2019.
- 609 Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In Yoshua Bengio
 610 and Yann LeCun (eds.), *3rd International Conference on Learning Representations, ICLR 2015,
 611 San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015.
- 612 Alex Krizhevsky and Geoffrey Hinton. Learning multiple layers of features from tiny images. 2009.
- 613 Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu.
 614 Towards deep learning models resistant to adversarial attacks. *arXiv preprint arXiv:1706.06083*,
 615 2017.
- 616 Shishira R. Maiya, Max Ehrlich, Vatsal Agarwal, Ser-Nam Lim, Tom Goldstein, and Abhinav
 617 Shrivastava. A frequency perspective of adversarial robustness. *CoRR*, abs/2111.00861, 2021.
- 618 AKM Iqtidar Newaz, Nur Imtiazul Haque, Amit Kumar Sikder, Mohammad Ashiqur Rahman, and
 619 A Selcuk Uluagac. Adversarial attacks to machine learning-based smart healthcare systems. In
 620 *GLOBECOM 2020-2020 IEEE Global Communications Conference*, pp. 1–6. IEEE, 2020.
- 621 Tianyu Pang, Kun Xu, Chao Du, Ning Chen, and Jun Zhu. Improving adversarial robustness via
 622 promoting ensemble diversity. In *International Conference on Machine Learning*, pp. 4970–4979.
 623 PMLR, 2019.
- 624 Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito,
 625 Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. Automatic differentiation in
 626 pytorch. 2017.
- 627 Rahul Rade and Seyed-Mohsen Moosavi-Dezfooli. Helper-based adversarial training: Reducing
 628 excessive margin to achieve a better accuracy vs. robustness trade-off. In *ICML 2021 Workshop on
 629 Adversarial Machine Learning*, 2021.
- 630 Edward Raff, Jared Sylvester, Steven Forsyth, and Mark McLean. Barrage of random transforms for
 631 adversarially robust defense. In *Proceedings of the IEEE/CVF Conference on Computer Vision
 632 and Pattern Recognition*, pp. 6528–6537, 2019.
- 633 Giulio Rossolini, Federico Nesti, Gianluca D’Amico, Saasha Nair, Alessandro Biondi, and Giorgio
 634 Buttazzo. On the real-world adversarial robustness of real-time semantic segmentation models
 635 for autonomous driving. *IEEE Transactions on Neural Networks and Learning Systems*, pp. 1–15,
 636 2023. doi: 10.1109/TNNLS.2023.3314512.
- 637 Andrei A. Rusu, Dan Andrei Calian, Sven Gowal, and Raia Hadsell. Hindering adversarial attacks
 638 with implicit neural representations. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba
 639 Szepesvári, Gang Niu, and Sivan Sabato (eds.), *International Conference on Machine Learning,
 640 ICML 2022, 17-23 July 2022, Baltimore, Maryland, USA*, volume 162 of *Proceedings of Machine
 641 Learning Research*, pp. 18910–18934. PMLR, 2022.

- 648 Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image
 649 recognition. *arXiv preprint arXiv:1409.1556*, 2014.
 650
- 651 James C Spall. Multivariate stochastic approximation using a simultaneous perturbation gradient
 652 approximation. *IEEE transactions on automatic control*, 37(3):332–341, 1992.
 653
- 654 Jacob Springer, Melanie Mitchell, and Garrett Kenyon. A little robustness goes a long way: Lever-
 655 aging robust features for targeted transfer attacks. *Advances in Neural Information Processing
 656 Systems*, 34:9759–9773, 2021.
- 657 Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow,
 658 and Rob Fergus. Intriguing properties of neural networks. In *2nd International Conference on
 659 Learning Representations, ICLR 2014*, 2014.
 660
- 661 Florian Tramèr, Alexey Kurakin, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick Mc-
 662 Daniel. Ensemble adversarial training: Attacks and defenses. In *International Conference on
 663 Learning Representations*, 2018.
- 664 Dimitris Tsipras, Shibani Santurkar, Logan Engstrom, Alexander Turner, and Aleksander Madry. Ro-
 665 bustness may be at odds with accuracy. In *International Conference on Learning Representations*,
 666 2018.
 667
- 668 Zekai Wang, Tianyu Pang, Chao Du, Min Lin, Weiwei Liu, and Shuicheng Yan. Better diffusion
 669 models further improve adversarial training. In Andreas Krause, Emma Brunskill, Kyunghyun
 670 Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett (eds.), *International Conference
 671 on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, volume 202 of
 672 *Proceedings of Machine Learning Research*, pp. 36246–36263. PMLR, 2023a.
 673
- 674 Zekai Wang, Tianyu Pang, Chao Du, Min Lin, Weiwei Liu, and Shuicheng Yan. Better diffusion
 675 models further improve adversarial training. In *International Conference on Machine Learning*,
 676 pp. 36246–36263. PMLR, 2023b.
 677
- 678 Zifan Wang, Yilin Yang, Ankit Shrivastava, Varun Rawal, and Zihao Ding. Towards frequency-based
 679 explanation for robust cnn. *arXiv preprint arXiv:2005.03141*, 2020.
 680
- 681 Futa Waseda, Sosuke Nishikawa, Trung-Nghia Le, Huy H Nguyen, and Isao Echizen. Closer
 682 look at the transferability of adversarial examples: How they fool different models differently.
 683 In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pp.
 1360–1368, 2023.
 684
- 685 Sven-Ake Wegner. Lecture notes on high-dimensional data. *arXiv preprint arXiv:2101.05841*, 2021.
 686
- 687 Yuancheng Xu, Yanchao Sun, Micah Goldblum, Tom Goldstein, and Furong Huang. Exploring and
 688 exploiting decision boundary dynamics for adversarial robustness. In *The Eleventh International
 689 Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*, 2023.
 690
- 691 Huanrui Yang, Jingyang Zhang, Hongliang Dong, Nathan Inkawich, Andrew Gardner, Andrew
 692 Touchet, Wesley Wilkes, Heath Berry, and Hai Li. Dverge: diversifying vulnerabilities for
 693 enhanced robust generation of ensembles. *Advances in Neural Information Processing Systems*,
 33:5505–5515, 2020.
 694
- 695 Zhuolin Yang, Linyi Li, Xiaojun Xu, Shiliang Zuo, Qian Chen, Pan Zhou, Benjamin Rubinstein,
 Ce Zhang, and Bo Li. Trs: Transferability reduced ensemble via promoting gradient diversity and
 696 model smoothness. *Advances in Neural Information Processing Systems*, 34:17642–17655, 2021.
 697
- 698 Mehmet Kerim Yucel, Ramazan Gokberk Cinbis, and Pinar Duygulu. Hybridaugment++: Unified fre-
 699 quency spectra perturbations for model robustness. In *Proceedings of the IEEE/CVF International
 700 Conference on Computer Vision*, pp. 5718–5728, 2023.
 701
- Sergey Zagoruyko and Nikos Komodakis. Wide residual networks. *CoRR*, abs/1605.07146, 2016.

702 Hongyang Zhang, Yaodong Yu, Jiantao Jiao, Eric P. Xing, Laurent El Ghaoui, and Michael I. Jordan.
703 Theoretically principled trade-off between robustness and accuracy. In Kamalika Chaudhuri
704 and Ruslan Salakhutdinov (eds.), *Proceedings of the 36th International Conference on Machine*
705 *Learning, ICML 2019, 9-15 June 2019, Long Beach, California, USA*, volume 97 of *Proceedings*
706 *of Machine Learning Research*, pp. 7472–7482. PMLR, 2019.

707 Wen Zhou, Xin Hou, Yongjun Chen, Mengyun Tang, Xiangqi Huang, Xiang Gan, and Yong Yang.
708 Transferable adversarial perturbations. In *Proceedings of the European Conference on Computer*
709 *Vision (ECCV)*, pp. 452–467, 2018.

711 Yi Zhu, Chenglin Miao, Tianhang Zheng, Foad Hajiaghajani, Lu Su, and Chunming Qiao. Can we
712 use arbitrary objects to attack lidar perception in autonomous driving? In *Proceedings of the 2021*
713 *ACM SIGSAC Conference on Computer and Communications Security*, pp. 1945–1960, 2021.

714

715

716

717

718

719

720

721

722

723

724

725

726

727

728

729

730

731

732

733

734

735

736

737

738

739

740

741

742

743

744

745

746

747

748

749

750

751

752

753

754

755

756 **A OMITTED PROOFS**

757 **A.1 PROOF OF INEQUALITY (6):**

$$\begin{aligned} & \sum_{y_t \in \mathcal{Y}} \text{Vr}(F_E, y_t) \\ &= \sum_{y_t \in \mathcal{Y}} \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\}] \\ &= \sum_{y_t \in \mathcal{Y}} \sum_{(x,y) \in \mathcal{D}} p_{(x,y)} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\}]. \end{aligned}$$

765 Then, we interchange the order of summation, and so the above equation is equal to

$$\begin{aligned} & \sum_{(x,y) \in \mathcal{D}} p_{(x,y)} \sum_{y_t \in \mathcal{Y}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\}] \\ &= \mathbb{E}_{(x,y) \sim \mathcal{D}} \left[\sum_{y_t \in \mathcal{Y}} \mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\} \right]. \end{aligned}$$

766 For each (x, y) , without loss of generality, let $F_E(\mathcal{A}(x)) = y_0$. For $y_t \neq y_0$, $\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\} = 0$. For $y_t = y_0$, $\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y_t\} = \mathbb{I}\{F_E(x) = y\}$. So the 767 above equation is equal to

$$\begin{aligned} & \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y\}] \\ &= \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge (F_E(\mathcal{A}(x)) \neq y \vee F_E(\mathcal{A}(x)) = y)\}] \\ &= \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) \neq y\} + \mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y\}]. \end{aligned}$$

768 We split $\mathbb{I}(\cdot)$ because $F_E(\mathcal{A}(x)) \neq y$ and $F_E(\mathcal{A}(x)) = y$ are mutually exclusive. Then, the above 769 equation is equal to

$$\begin{aligned} & \text{Vr}(F_E) + \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathbb{I}\{F_E(x) = y \wedge F_E(\mathcal{A}(x)) = y\}] \\ &\geq \text{Vr}(F_E) \end{aligned}$$

770 Overall, we can obtain the inequality (6): $\sum_{y_t \in \mathcal{Y}} \text{Vr}(F_E, y_t) \geq \text{Vr}(F_E)$

771 **B FREQUENCY SELECTION**

772 Figure 4 illustrates an example to show that, if we keep high-amplitude frequencies and remove 773 low-amplitude ones, the image is changed slightly even with adding certain noise (i.e., we can 774 still recognize the ground truth from the modified image). On the other hand, if we keep the low- 775 amplitude frequencies only, the semantic information is almost missing. This observation suggests 776 that high-amplitude frequency features are more strongly related to the semantic information of 777 image.

778 **C RANDOM NOISE BASED TRANSFORMATION**

779 **Random noise based transformation:** This approach substitutes the identified non-robust frequencies with Gaussian noise. For an $N \times N$ image, we take the non-robust frequencies based on the pre-specified threshold τ , and replace them with random vector for each sub-model in our experiment. In particular, to further increase the randomness, we perform this transformation for each epoch in the training stage. If we select the top s non-robust frequencies, the overall dimensionality of the edited random feature should be $s \times E$ (we concatenate those s -dimensional features together), where E is the number of epochs. For example, if $N = 32$, $s = N^2/2$, and $E = 200$, the overall dimensionality can be as large as 10^5 . Because these M features are random and have high dimensions, they are very likely to be nearly orthogonal with each other (this phenomenon in high-dimensional geometry can be proved by the central limit theorem (Wegner, 2021)). As a consequence, they tend to yield diverse training results for the sub-models.

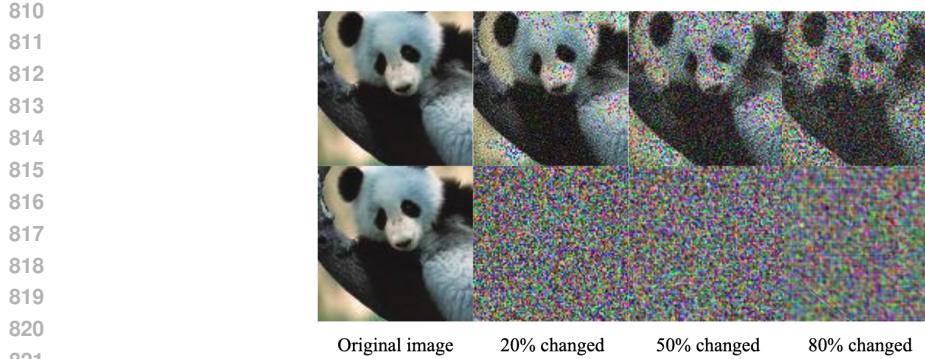


Figure 4: The first and second rows are the figures by adding random noise to high-amplitude and low-amplitude frequencies, respectively. “20% changed” for the first row means we remove the 20% lowest-amplitudes frequencies, and add small noise to the remaining high-amplitude frequencies. “20% changed” for the second row means we remove the 20% highest-amplitude frequencies, and add small noise to the remaining low-amplitude frequencies. “50% changed” and “80% changed” follow the same procedure as “20% changed”.

The implementation details are as follows. Given an image x , we perform Fourier Transform on x and also on a generated Gaussian noise n_0 . Then, we can obtain the low-amplitude frequencies and high-amplitude frequencies of x by setting an amplitude threshold. Next, we generate two masks (M_1 and M_2) to select high-amplitude frequencies and low-amplitude frequencies. We add the low-amplitude frequencies of n_0 (i.e., $M_2(n_0)$) to high-amplitude frequencies of x (i.e., $M_1(x)$), and obtain the transformation of x (denoted as $\pi(x)$). Finally, we transform $\pi(x)$ to time domain by inverse Fourier transform and train the model with $\pi(x)$.

D ALGORITHM OF FDT

Algorithm 1 shows the overall framework of training an ensemble model with FDT. It illustrates that our data transformation is performed at each iteration.

Algorithm 1 Training ensemble model with FDT

```

Input: dataset  $\mathcal{X} \times \mathcal{Y}$ , the number of sub-models  $M$ , and the epoch number  $E$ ,
Output: sub-model  $\beta_1, \beta_2, \dots, \beta_M$ 
for  $i = 1$  to  $E$  do
    for each  $(x, y)$  in  $\mathcal{X} \times \mathcal{Y}$  do
        Perform frequency selection on  $x$ ;
        for  $j = 1$  to  $M$  do
             $\pi_j(x) \leftarrow$  Perform data transformation on  $x$  at the selected frequencies;
            train  $\beta_j$  on  $(\pi_j(x), y)$ 
        end for
    end for
end for

```

Algorithm 2 shows the details of targeted-attack transformation method on the whole dataset. For each specific image x , we obtain the targeted class according to the allocation scheme mentioned in “Stage (1)”. Then, we use targeted PGD attack to obtain the adversarial sample x' . After that, we perform Fourier Transform on x and x' , and we can obtain the low-amplitude frequencies and high-amplitude frequencies of x by setting an amplitude threshold. Next, we generate two masks (M_1 and M_2) to select high-amplitude frequencies and low-amplitude frequencies. We add the low-amplitude frequencies of x' (i.e., $M_2(x')$) to high-amplitude frequencies of x (i.e., $M_1(x)$), and obtain the transformation of x (denoted as $\pi(x)$). Finally, we transform $\pi(x)$ to time domain by inverse Fourier transform.

864 **Algorithm 2 Targeted-attack Transformation**

865 **Input:** dataset P_{ori} , number M , steps s , class number k
 866 **Output:** Transformed data P_1, P_2, \dots, P_M
 867 Divide the dataset P_{ori} into k parts $\{C_1, C_2, \dots, C_k\}$ according to labels
 868 Randomly partition the dataset C_j equally into disjoint $k - 1$ parts $\{C_{j,1}, C_{j,2}, \dots, C_{j,k-1}\}$
 869 Initialize P_1, P_2, \dots, P_M with empty set;
 870 $m \leftarrow 0$
 871 **for** $j = 1$ to k **do**
 872 **for** $i = 1$ to $k - 1$ **do**
 873 $C'_{j,i} \leftarrow$ calculate targeted attack example in $C_{j,i}$ with label $i + j \mod k$ and perform data
 874 transformation on each image;
 875 **for** $s = 1$ to $\lceil \frac{M}{2} \rceil + 1$ **do**
 876 $m \leftarrow m + 1 \mod M$;
 877 Append $C'_{j,i}$ to P_m ;
 878 **end for**
 879 **end for**
 880 **end for**

882 E IMPLEMENTATION

883

884 In this section, we provide more experimental details. In our work, we utilize the CIFAR-10
 885 (Krizhevsky & Hinton, 2009), CIFAR-100 (Krizhevsky & Hinton, 2009), and Tiny-ImageNet-200
 886 (Deng et al., 2009). We acknowledge Alex Krizhevsky for the creation of CIFAR-10 and CIFAR-100,
 887 which are released under the MIT license. We also acknowledge Deng et al. (2009) for the ImageNet
 888 dataset Tiny-ImageNet-200, which permits non-commercial research use. In the testing process, the
 889 primary reason for selecting FGSM (Madry et al., 2017), PGD (Carlini et al., 2019), BIM (Goodfellow
 890 et al., 2015), MIM (Dong et al., 2018), CW(Carlini & Wagner, 2017) as attack methods is to keep
 891 consistent with the baseline methods from the literature. Further, we select AA (Croce & Hein,
 892 2020) because it is also a popular attack method and more powerful than those base methods. To
 893 reduce the computational complexity of targeted attacks, we leverage the transferability of adversarial
 894 examples and utilize a pre-trained simple network (VGG11(Simonyan & Zisserman, 2014)) structure
 895 for targeted attacks.

896 Further, we introduce the implement of “FDT-random”, “FDT-target” and “FDT-hybrid” here. For
 897 “FDT-random”, we perform Fourier Transform on x and also on a generated Gaussian noise n_0 . Then,
 898 we can obtain the low-amplitude frequencies and high-amplitude frequencies of x by setting an
 899 amplitude threshold. Next, we generate two masks (M_1 and M_2) to select high-amplitude frequencies
 900 and low-amplitude frequencies. We add the low-amplitude frequencies of n_0 (i.e., $M_2(n_0)$) to
 901 high-amplitude frequencies of x (i.e., $M_1(x)$), and obtain the transformation of x (denoted as $\pi(x)$).
 902 Finally, we transform $\pi(x)$ to time domain by inverse Fourier transform and train the model with
 903 $\pi(x)$. For “FDT-target”, we obtain the targeted class according to the allocation scheme mentioned
 904 in “Stage (1)”. Then, we use targeted PGD attack to obtain the adversarial sample x' . After that,
 905 perform the same steps as with FDT-random (we substitute n_0 with x'). For FDT-hybrid, we set two
 906 frequency selection thresholds τ_1 and τ_2 ($\tau_1 < \tau_2$), and generate three masks to select the frequencies:
 907 M_1 for the high-amplitude frequencies (amplitude $> \tau_2$), M_2 for the middle part ($\tau_1 <$ amplitude
 908 $< \tau_2$), and M_3 for the small part (amplitude $< \tau_1$). Next, we combine $M_1(x)$, $M_2(x')$ and $M_3(n_0)$
 909 to obtain the transformation $\pi(x)$.

910 F ADDITIONAL EXPERIMENTAL RESULTS

911

912 In this section, we provide more experimental results. Firstly, we evaluate the performance of FDT
 913 under black-box attacks on the CIFAR-10 and CIFAR-100 in Appendix F.1 Then we present the
 914 trade-off between clean accuracy and robust accuracy on the CIFAR-100 using FDT method in
 915 Appendix F.2. This trade-off sheds light on the effectiveness of FDT with changing the trade-off
 916 parameter. Additionally, in Appendix F.3, we compare the transferability across various sub-models
 917 with the baseline methods. Then, we extend our experiments to SVHN, Tiny-ImageNet-200, and
 WideResNet-28-10 in Appendix F.4. We also conduct the ablation studies on weakness set allocation,

amplitude-based selection and model architecture in Appendix F.5. Furthermore, we compare our method with more state-of-the-art methods in Appendix F.6.

F.1 RESULTS FOR BLACK-BOX ATTACK

In the black-box setting, the attacker's knowledge usually is limited to the original training dataset and has no information about the model. This setting represents a more practical attack scenario. The attacker can train a surrogate model to generate transferable adversarial examples and transfer them to the target ensemble model. We utilize a single ResNet-20 model as the surrogate model. Adversarial examples are generated on the surrogate model using the SPSA algorithm (Spall, 1992). Figure 5 shows the robust accuracy of ensemble models against black-box attacks under different degrees of perturbation. As we can see, FDT-hybrid ensemble training strategies outperform the other ensemble training strategy against black-box attacks both on CIFAR10 and CIFAR100. The results show the competitive robustness of our method against black-box attacks.

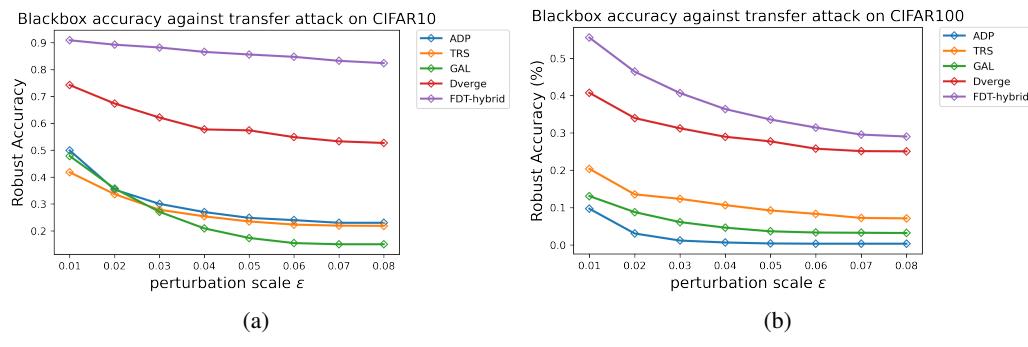


Figure 5: Robust Accuracy for different ensemble models against black-box attack with different perturbation scale ϵ .

F.2 TRADE-OFF BETWEEN CLEAN AND ROBUST ACCURACY

In this section, we explore the trade-off between clean accuracy and robust accuracy by varying the frequency selection threshold τ (as mentioned in Section 4.2). To assess the adversarial robustness, we utilize the PGD attack under l_∞ perturbations of size $\epsilon = 0.01$ as a benchmark. We train a set of ResNet-20 FDT-hybrid models on CIFAR-10 CIFAR-100 with various frequency selection threshold $\tau \in \{0.4, 0.6, 0.8, 1.0, 1.2, 1.6\}$. Figure 6 shows that the ensemble model has lower clean accuracy and higher robust accuracy with the increasing of τ .

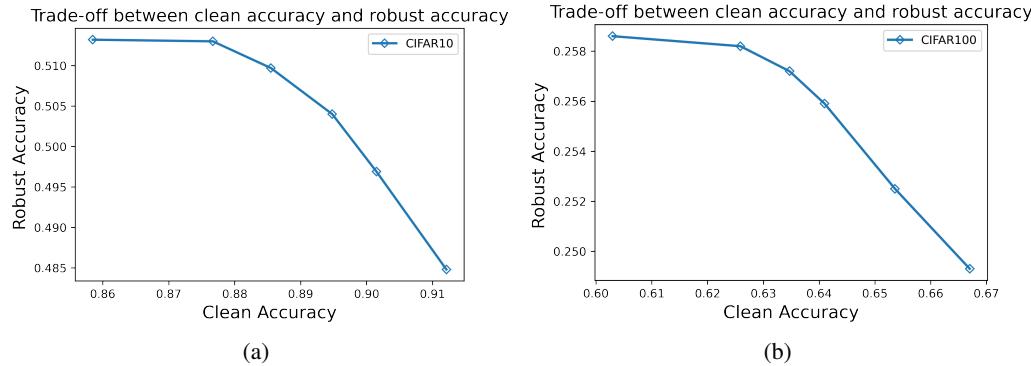
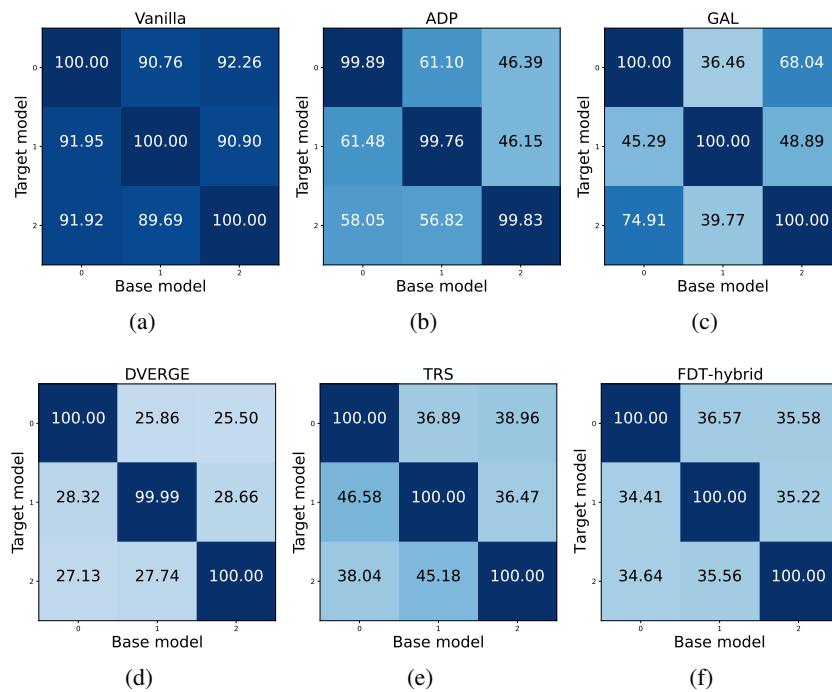


Figure 6: (a) shows the trade-off on CIFAR-10 while (b) on CIFAR-100. From left to right, we decrease the trade-off parameter τ for FDT.

972 F.3 TRANSFERABILITY ACROSS VARIOUS SUB-MODELS
 973

974 To further investigate the diversity between sub-models, we conduct an analysis by generating
 975 adversarial examples using one sub-model and evaluating their accuracy on other target sub-models.
 976 The accuracy of these adversarial examples among sub-models is visualized in Figure 7, considering
 977 different ensemble training methods on the CIFAR10 dataset. The experimental results indicate
 978 that FDT exhibits comparable performance to DVERGE and TRS in reducing the transferability of
 979 adversarial samples across different sub-models. This demonstrates that FDT not only enhances
 980 the diversity of weaknesses within the dataset but also weakens the transferability of adversarial
 981 examples between sub-models.



1005 Figure 7: Pair-wise adversarial transferability between sub-models against PGD attack with $\epsilon = 0.02$
 1006 on CIFAR-10.
 1007

1008 F.4 EXPERIMENTAL RESULTS ON ADDITIONAL DATASETS AND ARCHITECTURE
 1009

1010 In this section, we extend our experiments to additional datasets (SVHN, Tiny-ImageNet-200) and
 1011 architecture (WideResNet-28-10).

1012 Table 3 presents the performance of ensemble methods trained with ResNet-20 on SVHN against
 1013 several widely used white-box attacks. The experimental results demonstrate that all ensemble
 1014 models achieve comparable levels of clean accuracy. Specifically, the FDT approach exhibits better
 1015 robust accuracy than the other methods. These observations highlight the effectiveness of FDT in
 1016 achieving favorable clean accuracy and robustness of ensemble models, while concurrently mitigating
 1017 the computational cost during the training phase.

1018 We also extend our experiment to the sub-models trained with WideResNet-28-10 on CIFAR-10.
 1019 Table 4 shows the performance of the models facing various whitebox attacks. The results indicate
 1020 that FDT maintains good performance even on more complex network structures. We also evaluated
 1021 the robustness of an ensemble of eight sub-models, with the results presented in Table 5.

1022 Table 6 is the result of ensemble methods trained with WideResNet-28-10 on Tiny-ImageNet-200.
 1023 We test the robustness of different methods under widely used white-box attacks. Due to the high
 1024 time complexity of the TRS, we do not compare with it here. The experimental results show that
 1025 all ensemble models achieve comparable levels of clean accuracy while FDT-hybrid achieves better
 1026 robust accuracy than other methods.

1026
 1027 Table 3: Robust Accuracy (%) of different ensemble methods against white-box attacks on SVHN.
 1028 The ϵ and λ stand for the l_∞ norm of the adversarial perturbation and the coefficient of C&W attack
 1029 respectively. The last column is the ensemble model trained with FDT-hybrid.

SVHN	ADP	GAL	DVERGE	TRS	FDT-hybrid
clean accuracy	96.83	94.66	96.28	94.52	96.73 ± 0.12
FGSM ($\epsilon=0.01$)	84.38	80.2	85.6	72.87	90.13 ± 0.09
FGSM ($\epsilon=0.02$)	78.08	41.5	81.4	53.9	86.78 ± 0.07
PGD ($\epsilon=0.01$)	51.01	50.1	53.31	54.43	59.42 ± 0.07
PGD ($\epsilon=0.02$)	17.74	8.24	17.42	18.86	22.74 ± 0.04
BIM ($\epsilon=0.01$)	54.38	47.73	52.08	53.71	57.91 ± 0.08
BIM ($\epsilon=0.02$)	21.26	8.1	14.58	18.05	20.23 ± 0.05
MIM ($\epsilon=0.01$)	61.24	51.96	58.51	56.32	62.14 ± 0.08
MIM ($\epsilon=0.02$)	24.84	5.14	23.22	21.95	25.37 ± 0.04
AA ($\epsilon=0.01$)	49.92	48.39	52.02	52.83	57.54 ± 0.09
AA ($\epsilon=0.02$)	16.13	6.90	16.95	17.48	20.12 ± 0.05
C&W ($\lambda=0.1$)	55.81	49.94	66.82	52.74	72.14 ± 0.11

1044
 1045 Table 4: Robust Accuracy (%) of different ensemble methods against white-box attacks on CIFAR-10.
 1046 The ϵ and λ stand for the l_∞ norm of the adversarial perturbation and the coefficient of C&W attack
 1047 respectively. The architecture of sub-model is WRN-28-10.

CIFAR-10	ADP	GAL	DVERGE	FDT-hybrid
clean accuracy	92.99	82.14	94.32	94.18 ± 0.06
FGSM ($\epsilon=0.01$)	60.04	44.94	71.01	80.64 ± 0.05
FGSM ($\epsilon=0.02$)	51.69	36.83	50.43	60.09 ± 0.05
PGD ($\epsilon=0.01$)	11.09	22.10	44.25	64.64 ± 0.07
PGD ($\epsilon=0.02$)	2.54	5.06	13.27	26.0 ± 0.03
BIM ($\epsilon=0.01$)	15.81	22.62	46.53	67.36 ± 0.10
BIM ($\epsilon=0.02$)	4.50	5.43	17.38	32.36 ± 0.06
MIM ($\epsilon=0.01$)	18.18	25.97	44.21	64.36 ± 0.08
MIM ($\epsilon=0.02$)	4.72	7.81	12.83	25.64 ± 0.05
AA ($\epsilon=0.01$)	9.38	19.34	43.23	63.45 ± 0.08
AA ($\epsilon=0.02$)	1.17	3.93	12.49	25.23 ± 0.04
C&W ($\lambda=0.1$)	37.81	19.05	46.32	47.23 ± 0.10

F.5 MORE ABLATION STUDIES

1065 In this section, We explore the ablation studies on weakness set allocation, amplitude-based selection
 1066 and model architecture.

1067 Table 7 presents the results across different model architectures, including ResNet20, ResNet50,
 1068 WRN28-10, and WRN34-10. While larger models generally achieve higher clean and robust accuracy,
 1069 the results suggest that our method consistently enhances robustness under various attack scenarios,
 1070 demonstrating its applicability across diverse architectures.

1071 Table 8 compares the performance of FDT-hybrid with different weakness set allocation methods on
 1072 CIFAR-10. The results indicate that our proposed allocation method achieves better clean accuracy
 1073 and robustness under various attack scenarios than randomly uniform allocation.

1074 Table 9 presents the results of FDT-hybrid with various combinations of selection thresholds τ_1 and
 1075 τ_2 on CIFAR-10. The experiments reveal the impact of different thresholds on both clean accuracy
 1076 and robustness under adversarial attacks. As τ_2 increases, robustness improves across all metrics, but
 1077 clean accuracy decreases. For a fixed τ_2 , increasing τ_1 generally leads to a trade-off between clean
 1078 accuracy and robustness. Setting $\tau_1 = 0.2$ and $\tau_2 = 0.8$ achieves a relatively balanced performance,
 1079 maintaining both competitive clean accuracy and robust accuracy under various attacks.

1080
 1081 Table 5: Robust Accuracy (%) of an ensemble of eight sub-models against white-box attacks on
 1082 CIFAR-10. The ϵ and λ stand for the l_∞ norm of the adversarial perturbation and the coefficient of
 1083 C&W attack respectively. The architecture of sub-model is WRN-28-10.

	CIFAR-10	FDT-hybrid
clean accuracy	93.72 ± 0.11	
FGSM ($\epsilon=0.01$)	86.31 ± 0.07	
FGSM ($\epsilon=0.02$)	67.29 ± 0.06	
PGD ($\epsilon=0.01$)	72.02 ± 0.07	
PGD ($\epsilon=0.02$)	45.42 ± 0.05	
BIM ($\epsilon=0.01$)	73.68 ± 0.10	
BIM ($\epsilon=0.02$)	44.53 ± 0.06	
MIM ($\epsilon=0.01$)	71.36 ± 0.06	
MIM ($\epsilon=0.02$)	45.24 ± 0.06	
AA ($\epsilon=0.01$)	70.45 ± 0.08	
AA ($\epsilon=0.02$)	44.23 ± 0.07	
C&W ($\lambda=0.1$)	72.37 ± 0.11	

1099 Table 6: Robust Accuracy (%) of different ensemble methods against white-box attacks on Tiny-
 1100 ImageNet-200. The ϵ and λ stand for the l_∞ norm of the adversarial perturbation and the coefficient
 1101 of C&W attack respectively. The last column is the ensemble model trained with FDT-hybrid.

Tiny-ImageNet-200	ADP	GAL	DVERGE	FDT-hybrid
clean accuracy	49.88	45.7	51.46	64.21 ± 0.06
FGSM ($\epsilon=0.01$)	10.46	1.24	22.82	21.73 ± 0.04
FGSM ($\epsilon=0.02$)	4.38	0.59	18.42	19.28 ± 0.04
PGD ($\epsilon=0.01$)	0.02	0.02	3.6	4.76 ± 0.02
PGD ($\epsilon=0.02$)	0.02	0.01	0.34	0.45 ± 0.01
BIM ($\epsilon=0.01$)	0.07	0.02	3.35	4.81 ± 0.03
BIM ($\epsilon=0.02$)	0.03	0.01	0.28	0.32 ± 0.00
MIM ($\epsilon=0.01$)	0.11	0.02	4.36	6.13 ± 0.03
MIM ($\epsilon=0.02$)	0.03	0.01	0.41	0.48 ± 0.00
AA ($\epsilon=0.01$)	0	0	0	2.66 ± 0.02
AA ($\epsilon=0.02$)	0	0	0	0.02 ± 0.00
CW ($\lambda=0.01$)	2.36	0.13	9.54	19.47 ± 0.06

F.6 COMPARE WITH ADVERSARIAL TRAINING

1116 We use target attacks in our data transformation, which differs significantly from adversarial training.
 1117 First, we employ a simple pre-trained network (VGG11 in our experiment) to compute adversarial
 1118 examples, thereby accelerating the training process. Second, we only utilize the low amplitude
 1119 part of the adversarial examples for data transformation, which helps maintain the model’s clean
 1120 accuracy. We compare our method with several popular approaches (Wang et al., 2023b; Rade &
 1121 Moosavi-Dezfooli, 2021; Xu et al., 2023) on CIFAR-10 using AutoAttack under l_∞ perturbations
 1122 ($\epsilon = 8/255$). Wang et al. (2023b) generated training datasets using a diffusion model, followed by
 1123 adversarial training on these datasets. For fairness, we compare our method with the version proposed
 1124 by Wang et al. (2023b) that uses 50k generated images. Rade & Moosavi-Dezfooli (2021) used
 1125 “helper example” to help the adversarial training. Xu et al. (2023) proposed Dynamics-Aware Robust
 1126 Training, which encourages the decision boundary to adjust in a way that prioritizes increasing smaller
 1127 margins. We use WideResNet-28-10 as the sub-model and ensemble eight sub-models without using
 1128 generated data. The results in Table 10 indicate that, although the robustness of our method is not
 1129 the highest, it maintains clean accuracy with almost no decline. Moreover, our method does not
 1130 require additional generated data or adversarial training, and even with the need for ensembling,
 1131 the training efficiency remains relatively high. This suggests a potential way to enhance robustness while
 1132 minimizing the decrease in clean accuracy.

1134
 1135
 1136
 1137
 1138

Table 7: Robust Accuracy (%) of different ensemble methods against white-box attacks on Cifar10. The ϵ and λ stand for the l_∞ norm of the adversarial perturbation and the coefficient of C&W attack respectively.

CIFAR10	ResNet20	ResNet50	WRN28-10	WRN34-10
clean accuracy	90.02	93.23	94.18	94.63
FGSM ($\epsilon = 0.01$)	72.24	76.65	80.64	81.04
FGSM ($\epsilon = 0.02$)	58.04	58.59	60.09	60.92
PGD ($\epsilon = 0.01$)	48.48	60.23	64.64	65.38
PGD ($\epsilon = 0.02$)	20.01	24.35	26.00	27.42
BIM ($\epsilon = 0.01$)	48.57	60.43	67.36	68.29
BIM ($\epsilon = 0.02$)	16.63	23.57	32.36	33.86
MIM ($\epsilon = 0.01$)	51.48	60.81	64.36	64.71
MIN ($\epsilon = 0.02$)	20.09	24.54	25.64	26.42
AA ($\epsilon = 0.01$)	51.56	60.48	63.45	64.01
AA ($\epsilon = 0.02$)	19.42	24.21	25.23	26.39
CW ($\lambda = 0.01$)	56.08	56.55	57.23	57.52

1139
 1140
 1141
 1142
 1143
 1144
 1145
 1146
 1147
 1148
 1149
 1150
 1151
 1152
 1153
 1154

Table 8: Performance of FDT-hybrid with different weakness set allocation method on CIFAR-10. The other settings are consistent with those in Table 1.

allocation method	Clean accuracy	FGSM ($\epsilon=0.02$)	PGD ($\epsilon=0.02$)	AutoAttack ($\epsilon=0.02$)
Uniform Random	89.32	56.20	18.24	17.89
Ours	90.20	58.04	20.01	19.42

1161
 1162
 1163

Table 9: Performance of FDT-hybrid with different selection thresholds τ_1 and τ_2 on CIFAR-10. The other settings are consistent with those in Table 1.

thresholds	Clean accuracy	FGSM ($\epsilon=0.02$)	PGD ($\epsilon=0.02$)	AutoAttack ($\epsilon=0.02$)
$\tau_1 = 0.2, \tau_2 = 0.7$	91.03	56.62	17.74	17.60
$\tau_1 = 0.2, \tau_2 = 0.8$	90.20	58.04	20.01	19.42
$\tau_1 = 0.2, \tau_2 = 0.9$	89.46	58.48	20.12	19.57
$\tau_1 = 0.4, \tau_2 = 0.7$	89.75	56.89	17.93	17.82
$\tau_1 = 0.4, \tau_2 = 0.8$	89.08	58.21	20.01	19.47
$\tau_1 = 0.4, \tau_2 = 0.9$	88.44	58.53	20.09	19.61
$\tau_1 = 0.6, \tau_2 = 0.7$	89.62	53.35	15.27	14.63
$\tau_1 = 0.6, \tau_2 = 0.8$	88.84	55.33	15.42	15.24
$\tau_1 = 0.6, \tau_2 = 0.9$	88.12	55.46	15.83	15.47

1176
 1177
 1178
 1179

Table 10: Clean accuracy and robust accuracy (%) of different methods against AutoAttack under l_∞ perturbations ($\epsilon = 8/255$) on CIFAR-10.

CIFAR-10	clean accuracy	robust accuracy
(Wang et al., 2023b)	86.15	55.71
(Rade & Moosavi-Dezfooli, 2021)	84.90	49.08
(Xu et al., 2023)	85.55	54.69
OURS (FDT-hybrid)	93.72	34.61

1183
 1184
 1185
 1186
 1187